

Class 7: Machine Learning I

Audrey Nguyen

In this class we will explore clustering and dimensionality reduction methods.

##K-means

Make up some input data where we know what the answer should be.

```
tmp <- c(rnorm(30, -3), rnorm(30, +3))
x <- cbind(x=tmp, y=rev(tmp))
x
```

	x	y
[1,]	-1.8528805	2.7783032
[2,]	-2.9269715	2.5506729
[3,]	-2.8032023	5.8059108
[4,]	-1.1424205	2.8209750
[5,]	-3.0843831	1.9859521
[6,]	-2.5135990	1.6171640
[7,]	-3.4733960	1.8619089
[8,]	-2.4456156	4.0791860
[9,]	-3.8008968	1.6618588
[10,]	-2.2051226	3.6959530
[11,]	-4.1405486	3.9055803
[12,]	-3.6978199	2.1948239
[13,]	-3.6109422	3.6094984
[14,]	-2.8417313	3.6791965
[15,]	-2.7114965	1.6806124
[16,]	-3.7585458	4.0160263
[17,]	-2.0322426	0.3978394
[18,]	-3.6192585	2.9452341
[19,]	-2.8804125	2.2707954
[20,]	-2.7260425	1.3306288
[21,]	-2.5583816	2.3755967


```
[22,] -0.1245820  1.0559112
[23,] -2.7457598  3.6009875
[24,] -2.3835482  4.0643486
[25,] -2.3811922  3.7782588
[26,] -3.0613761  3.8827984
[27,] -3.0016676  4.5327120
[28,] -3.1104324  3.5029603
[29,] -3.1372732  3.0418920
[30,] -2.4555094  2.3412013
[31,]  2.3412013 -2.4555094
[32,]  3.0418920 -3.1372732
[33,]  3.5029603 -3.1104324
[34,]  4.5327120 -3.0016676
[35,]  3.8827984 -3.0613761
[36,]  3.7782588 -2.3811922
[37,]  4.0643486 -2.3835482
[38,]  3.6009875 -2.7457598
[39,]  1.0559112 -0.1245820
[40,]  2.3755967 -2.5583816
[41,]  1.3306288 -2.7260425
[42,]  2.2707954 -2.8804125
[43,]  2.9452341 -3.6192585
[44,]  0.3978394 -2.0322426
[45,]  4.0160263 -3.7585458
[46,]  1.6806124 -2.7114965
[47,]  3.6791965 -2.8417313
[48,]  3.6094984 -3.6109422
[49,]  2.1948239 -3.6978199
[50,]  3.9055803 -4.1405486
[51,]  3.6959530 -2.2051226
[52,]  1.6618588 -3.8008968
[53,]  4.0791860 -2.4456156
[54,]  1.8619089 -3.4733960
[55,]  1.6171640 -2.5135990
[56,]  1.9859521 -3.0843831
[57,]  2.8209750 -1.1424205
[58,]  5.8059108 -2.8032023
[59,]  2.5506729 -2.9269715
[60,]  2.7783032 -1.8528805
```

```
head(x)
```

	x	y
[1,]	-1.852881	2.778303
[2,]	-2.926972	2.550673
[3,]	-2.803202	5.805911
[4,]	-1.142420	2.820975
[5,]	-3.084383	1.985952
[6,]	-2.513599	1.617164

Quick plot of x to see the two graphs at -3, 3 and +3, -3

```
plot(x)
```



Use the `kmeans()` function setting `k` to 2 and `nstart = 20`

```
km <- kmeans(x, centers = 2, nstart = 20)
km
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	-2.774242	2.902160
2	2.902160	-2.774242

Clustering vector:

[illegible]

```
[1] 60.00383 60.00383
      (between_SS / total_SS =  89.0 %)
```

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```


km\$size

Q. What ‘component’ of your result details - cluster assignment/membership? - cluster center?

[illegible]


	x	y
1	-2.774242	2.902160
2	2.902160	-2.774242

```
plot(x, col = km$cluster)
points(km$centers, col = "blue", pch = 15, cex = 2)
```



Play with kmeans and ask for different number of clusters

```
km <- kmeans(x, centers = 4, nstart = 20)
plot(x, col = km$cluster)
points(km$centers, col = "blue", pch = 16, cex = 2)
```



Hierarchical Clustering

This is another very useful and widely employed clustering method which has the advantage over k-means in that it can help reveal the something of the true grouping in your data.

The `hclust()` function wants a distance matrix as input. We can get this from the `dist()` function.

```
d <- dist(x)
hc <- hclust(d)
hc
```


Call:
hclust(d = d)

Cluster method : complete
Distance : euclidean
Number of objects: 60

There is a plot message for hclust results:

```
plot(hc)
abline(h = 10, col = "red")
```

Cluster Dendrogram



To get my cluster membership vector, I need to “cut” my tree to yield sub-trees or branches with all the members of a given cluster residing on the same cut branch. The function to do this is `cutree()`.

```
groups <- cutree(hc, h = 10)
groups
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
[39] 1 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

It is often helpful to use the `k` = argument rather than the `h` = height of cutting with `cutree()`. This will cut the tree to yield the number of clusters you want.

```
cutree(hc, k = 4)
```

```
[1] 1 1 2 1 1 1 1 2 1 2 2 1 2 2 1 2 1 2 1 1 1 1 2 2 2 2 2 2 2 1 3 4 4 4 4 4 4 4
[39] 1 3 3 3 4 3 4 3 4 4 3 4 4 3 4 3 3 3 3 4 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3
```

#Principal Component Analysis (PCA)

The base R function for PCA is called `prcomp()`. Let's play with some 17D data (a very small dataset) and see how PCA can help.

#PCA of UK food data

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
head(x)
```

	X	England	Wales	Scotland	N.Ireland
1	Cheese	105	103	103	66
2	Carcass_meat	245	227	242	267
3	Other_meat	685	803	750	586
4	Fish	147	160	122	93
5	Fats_and_oils	193	235	184	209
6	Sugars	156	175	147	139

Q1. How many rows and columns are in your new data frame named `x`? What R functions could you use to answer these questions?

You can use the `dim()` function to return the number of rows and columns or `nrow()` to get rows and `ncol()` to get columns separately.

```
dim(x)
```

```
[1] 17 5
```

```
nrow(x)
```

```
[1] 17
```

```
ncol(x)
```

```
[1] 5
```

```
## Preview the first 6 rows  
head(x)
```

	X	England	Wales	Scotland	N.Ireland
1	Cheese	105	103	103	66
2	Carcass_meat	245	227	242	267
3	Other_meat	685	803	750	586
4	Fish	147	160	122	93
5	Fats_and_oils	193	235	184	209
6	Sugars	156	175	147	139

It appears that the row-names are incorrectly set as the first column of our `x` data frame (rather than set as proper row-names). We want 4 columns for the 4 countries instead. We can fix this with the function `rownames()` to the first column and then remove the troublesome first column (with the `-1` column index):

```
# Note how the minus indexing works  
rownames(x) <- x[, 1]  
x <- x[, -1]  
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139


```
# checking the dimensions again
dim(x)
```

```
[1] 17  4
```

An alternative approach to setting the correct row-names in this case would be to read the data file again and set the `row.names` argument of `read.csv()` to be the 1st column.


```
x <- read.csv(url, row.names = 1)
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

I prefer the 2nd approach, because it’s more concise. The 2nd approach is more robust, because it can be run multiple times without messing up the dimensions.


```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



Q3. Changing what optional argument in the above `barplot()` function results in the following plot (bars are stacked on each other)?


Setting `beside=FALSE` in the `barplot()` code would stack the bars.

```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

```
pairs(x, col = rainbow(10), pch = 16)
```



You can compare the countries from switching the axes based on which pair you're looking at.

Given points on the diagonal means that they're the same value as the other food categories from other countries.

Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

Comparing to other countries, there are less points on the diagonal, which means that N. Ireland has more distinct values for the food categories.

```
# Use the prcomp() PCA function
pca <- prcomp(t(x))
summary(pca)
```

Importance of components:


	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	4.189e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

A “PCA plot” (a.k.a “Score plot”, PC1vsPC2 plot, etc.)

```
pca$x
```


	PC1	PC2	PC3	PC4
England	-144.99315	2.532999	-105.768945	2.842865e-14
Wales	-240.52915	224.646925	56.475555	7.804382e-13
Scotland	-91.86934	-286.081786	44.415495	-9.614462e-13
N.Ireland	477.39164	58.901862	4.877895	1.448078e-13

```
plot(pca$x[, 1], pca$x[, 2], col = c("orange", "red", "blue", "darkgreen"), pch = 15)
```




Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.

```
# Plot PC1 vs PC2
plot(pca$x[, 1], pca$x[, 2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(x))
```



Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.

```
plot(pca$x[, 1], pca$x[, 2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], col = c("orange", "red", "blue", "darkgreen"), colnames(x))
```



Below we can use the square of `pca$sdev`, which stands for “standard deviation”, to calculate how much variation in the original data each PC accounts for.

```
v <- round( pca$sdev^2/sum(pca$sdev^2) * 100 )
v
```


```
[1] 67 29 4 0
```

```
## or the second row here...
z <- summary(pca)
z$importance
```

	PC1	PC2	PC3	PC4
Standard deviation	324.15019	212.74780	73.87622	4.188568e-14
Proportion of Variance	0.67444	0.29052	0.03503	0.000000e+00
Cumulative Proportion	0.67444	0.96497	1.00000	1.000000e+00

This information can be summarized in a plot of the variances (eigenvalues) with respect to the principal component number (eigenvector number), which is given below.

```
barplot(v, xlab="Principal Component", ylab="Percent Variation")
```



```
## Lets focus on PC1 as it accounts for > 90% of variance
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,1], las=2 )
```

